

Socket programming with TCP/UDP

Sockets

- a way to speak to other programs using standard Unix file descriptors (like a pipe)
- low level communication endpoint used for processing info across a network

Types of sockets:

1. Stream Sockets (*SOCK_STREAM*)
 - reliable 2 way connected communication streams
 - uses TCP (Transmission Control Protocol)
2. Datagram Sockets (*SOCK_DGRAM*)
 - unreliable, connectionless (although they can be *connect()*'d)
 - uses UDP (User Datagram Protocol)

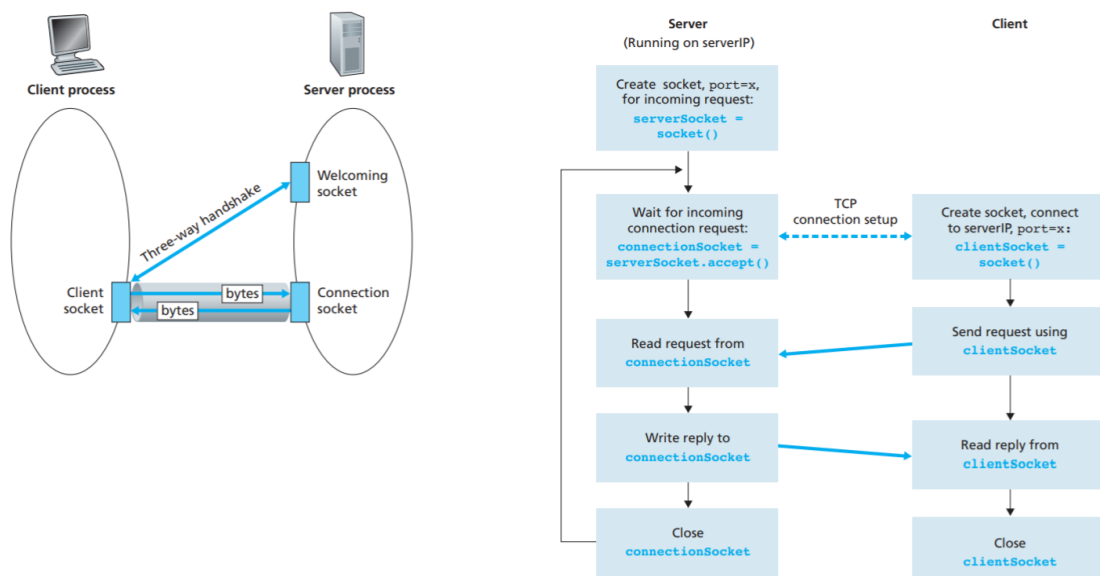
TCP API and client-server paradigm

TCP is a connection-oriented protocol → before the exchange of data a TCP connection must be established, one end to the client socket and the other to the server socket.

The connection is associated to the client and server socket addresses (IP + port number). Now, the data exchange happens by dropping the data into the TCP connection via its socket.

Client must contact server so:

- the server process must be running
- the server must have a special door (socket) that welcomes the client's contact
- client creates a local socket specifying the IP address and port number for the server
- the server accepts the contact and creates a new socket for that client's demand



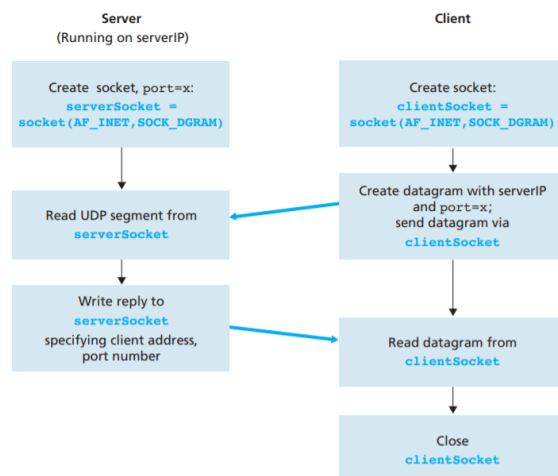
UDP API and client-server paradigm

UDP is a connection-less protocol → datagram oriented (individual messages).

We have no guarantee for datagram delivery and ordering.

Client must contact server so:

- the server process must be running
- client attaches to the packet a destination address (destination host's IP address and the destination socket's port number)
- the client's source address is also attached to the packet, but this process is done by the underlying operating system



Advantages: -less latency, higher bandwidth, broadcast (1 sender multiple receivers).

fdf

Programming

1. Structs and data handling

a) Socket descriptor → regular int

b) struct sockaddr → holds socket address info

```
struct sockaddr{
    unsigned short sa_family; //address family, AF_XXX
    char sa_data[14]; //14 bytes of protocol address
};
```

- sa_family will be AF_INET

▼ What is AF_INET?

Describes the type of addresses that our socket can communicate with, in this case Internet Protocol v4 addresses (IPv4)

- sa_data contains a destination address and port number for the socket

c) struct sockaddr_in → basically sockaddr but easier to manage

```
struct sockaddr_in {
    short int sin_family; //address family
    unsigned short int sin_port; //port nr
```

```
struct in_addr sin_addr; //internet address
unsigned char sin_zero[8]; //same size as struct sockaddr
}
```

- sin_zero should be set to all zeros with the function memset()
- sockaddr_in ↔ sockaddr so even tho socket() wants a sockaddr* you can give a sockaddr_in (cast it)

2. Conversions

Two types that you can convert: short (two bytes) and long (four bytes). Functions:

- htons() -- "Host to Network Short"
- htonl() -- "Host to Network Long"
- ntohs() -- "Network to Host Short"
- ntohl() -- "Network to Host Long"

!!! put your bytes in Network Byte Order before you put them on the network.

The sin_addr and sin_port need to be in Network Byte Order because they get encapsulated. So sin_family doesn't need to because it does not get sent out on the network.

3. IP addresses & how to deal with them

- inet_addr() → converts an IP address in numbers-and-dots notation into an unsigned long
- inet_ntoa() → reverse for inet_addr()

```
struct sockaddr_in ina;
ina.sin_addr.s_addr = inet_addr("10.12.110.57");
```

4. System calls & functions

▼ socket()

```
#include <sys/types.h>
#include <sys/socket.h>

int socket(int domain, int type, int protocol);
```

- domain → AF_INET
- type → what kind of socket it is (SOCK_STREAM/SOCK_DGRAM)
- protocol → set to '0' so the function will choose the correct protocol based on the type
- returns -1 on error

Socket options

```
int setsockopt(int sockfd, int level, int optname,
               const void *optval, socklen_t optlen);
```

Optname:

- SO_REUSEADDR – reuse local addresses – use it to get rid of the address already in use after you terminate your TCP/IP app abruptly
- SO_BROADCAST – enables broadcast = one sender – Universal(all) receiver(s)!

▼ bind()

= associate that socket with a port on your local machine

```
#include <sys/types.h>
#include <sys/socket.h>

int bind(int sockfd, struct sockaddr *my_addr, int addrlen);
```

- sockfd → the socket file descriptor
- my_addr → pointer to a struct sockaddr that contains info
- addrlen → set to sizeof(struct sockaddr)
- returns -1 on error
- ports below 1024 are reserved, so choose any port right up to 65535

Some of the process of getting your own IP address and/or port can be automated:

```
my_addr.sin_port = 0; //choose an unused port at random
my_addr.sin_addr.s_addr = INADDR_ANY; //use my IP address
//no need to put INADDR_ANY in network byte order bcs is actually 0
```

▼ connect()

= connect to a remote host

```
#include <sys/types.h>
#include <sys/socket.h>

int connect(int sockfd, struct sockaddr *serv_addr, int addrlen);
```

▼ listen()

= you wait for incoming connections → 2 steps: listen() and accept()

```
int listen(int sockfd, int backlog);
```

- backlog → nr of connections allowed on the incoming queue

If you're gonna listen, you'll make the next sequence of system calls: socket(), bind(), listen(), accept();

▼ accept()

- returns a brand new socket file descriptor

```
int accept(int sockfd, void *addr, int *addrlen);
```

▼ send() and recv()

```
int send(int sockfd, const void *msg, int len, int flags);
```

- returns the number of bytes actually sent out

```
int recv(int sockfd, void *buf, int len, unsigned int flags);
```

- returns the number of bytes actually read into the buffer

▼ sendto() and recvfrom()

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags,
               const struct sockaddr *dest_addr, socklen_t addrlen);
```

- returns nr. of bytes transmitted or -1 on error
- data size must fit < 64Kb
- 1 sendto must be consumed by exactly 1 recvfrom on the receiver

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags,
                 struct sockaddr *src_addr, socklen_t *addrlen);
```

- returns nr. of bytes received or -1 on error
- if not all data is read in one call the remaining is discarded
- cannot receive more than max 64Kb

▼ close() and shutdown()

- close the connection on your socket descriptor → close(sockfd);
- shutdown() → like close but more control; ret 0 on succes, -1 otherwise

```
int shutdown(int sockfd, int how);
//how can be : 0 - no more receives; 1 - no more sends; 2 - no more
```

▼ getpeername() — Who are you?

- The function getpeername() will tell you who is at the other end of a connected stream socket. Once you have their address, you can use inet_ntoa() or gethostbyaddr() to print or get more information. The synopsis:

```
#include <sys/socket.h>

int getpeername(int sockfd, struct sockaddr *addr, int *addrlen);
```

▼ gethostname() — Who am I?

- It returns the name of the computer that your program is running on. The name can then be used by gethostbyname(), below, to determine the IP address of your local machine.

```
#include <unistd.h>

int gethostname(char *hostname, size_t size);
```

▼ DNS

- DNS - Domain Name Service
- you tell it what the human-readable address is for a site, and it'll give you the IP address

```
#include <netdb.h>

struct hostent *gethostbyname(const char *name);
```

```

struct hostent {
    char    *h_name;
    char    **h_aliases;
    int     h_addrtype;
    int     h_length;
    char    **h_addr_list;
};
#define h_addr h_addr_list[0]

```

- `h_name` -- Official name of the host.
- `h_aliases` -- A NULL-terminated array of alternate names for the host.
- `h_addrtype` -- The type of address being returned; usually `AF_INET`.
- `h_length` -- The length of the address in bytes.
- `h_addr_list` -- A zero-terminated array of network addresses for the host. Host addresses are in Network Byte Order.
- `h_addr` -- The first address in `h_addr_list`.

▼ Example

```

/*
** getip.c -- a hostname lookup demo
*/

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <netdb.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

int main(int argc, char *argv[])
{
    struct hostent *h;

    if (argc != 2) { // error check the command line
        fprintf(stderr, "usage: getip address\n");
        exit(1);
    }

    if ((h=gethostbyname(argv[1])) == NULL) { // get the host info
        perror("gethostbyname");
        exit(1);
    }

    printf("Host name   : %s\n", h->h_name);
    printf("IP Address  : %s\n", inet_ntoa(*(struct in_addr *)h->h_addr));

    return 0;
}

```

5. Example

▼ `client.py` ← TCP

```

from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence)
modifiedSentence = clientSocket.recv(1024)
print 'From Server:', modifiedSentence
clientSocket.close()

```

▼ server.py ← TCP

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET, SOCK_STREAM)
serverSocket.bind(('', serverPort))
serverSocket.listen(1)
print 'The server is ready to receive'
while 1:
    connectionSocket, addr = serverSocket.accept()
    sentence = connectionSocket.recv(1024)
    capitalizedSentence = sentence.upper()
    connectionSocket.send(capitalizedSentence)
    connectionSocket.close()
```

▼ client.py ← UDP

```
from socket import *
serverName = 'hostname' #IP address or name of the server
serverPort = 12000
clientSocket = socket(socket.AF_INET, socket.SOCK_DGRAM)
message = raw_input('Input lowercase sentence:')
clientSocket.sendto(message, (serverName, serverPort)) #we attach the dest address to our message
modifiedMessage, serverAddress = clientSocket.recvfrom(2048) #sv address contains the IP address and the port
print modifiedMessage
clientSocket.close()
```

▼ server.py ← UDP

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET, SOCK_DGRAM)
serverSocket.bind(('', serverPort))
print "The server is ready to receive"
while 1:
    message, clientAddress = serverSocket.recvfrom(2048)
    #we make use of client address because we have to know where to send our reply
    modifiedMessage = message.upper()
    serverSocket.sendto(modifiedMessage, clientAddress)
```

▼ client.c ← UDP

```
//UDP client in the internet domain

#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netdb.h>
#include <stdlib.h>
#include <string.h>
#include <sys/unistd.h>

void error(char* msg){
    perror(msg);
    exit(1);
}

int main(int argc, char **argv) {
    int sockfd, length, n;
    struct sockaddr_in server, from; //ip addressing (ip, port, type)
    struct hostent *hp; //DNS
    char buffer[1024];

    if(argc < 2){
        printf("Usage: %s <port> \n", argv[0]);
        exit(1);
    }
```

```

    }

    sockfd = socket(AF_INET, SOCK_DGRAM, 0);

    if(sockfd < 0)
        error("Error on opening client socket");

    server.sin_family = AF_INET;
    server.sin_addr.s_addr = inet_addr("192.168.100.40");
    server.sin_port = htons(atoi(argv[1]));

    length = sizeof(struct sockaddr_in);

    printf("Please enter the message: ");
    bzero(buffer, 1024);
    fgets(buffer, 1023, stdin);

    n = sendto(sockfd, buffer, strlen(buffer), 0, (struct sockaddr*)&server, length);

    if(n<0)
        error("Error on sending information to server");

    n = recvfrom(sockfd, buffer, 1023, 0, (struct sockaddr*)&from, &length);

    if(n<0)
        error("Error on receiving information from server");

    write(1, "Got an ack: ", 12);
    write(1, buffer, n);

    return 0;
}

```

▼ server.c ← UDP

```

/*
 * Create a datagram server, the port is passed
 * as an argument. The server runs forever.
 */

#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>
#include <unistd.h>
#include <stdlib.h>
#include <strings.h>
#include <string.h>

void error(char* msg){
    perror(msg);
    exit(1);
}

int main(int argc, char** argv) {
    int sockfd, length, fromlen, n;
    struct sockaddr_in server;
    struct sockaddr_in from;
    char buf[1024];

    if(argc < 2){
        fprintf(stderr, "Error, no port provided\n");
        exit(1);
    }

    sockfd = socket(AF_INET, SOCK_DGRAM, 0);

    if(sockfd<0)
        error("Error on opening server socket");

    length = sizeof(server);
    bzero(&server, length);
    server.sin_addr.s_addr = INADDR_ANY;
    server.sin_family = AF_INET;

```



```

server.sin_port = htons(atoi(argv[1]));

if(bind(sockfd, (struct sockaddr*)&server, length) < 0)
    error("Error on binding");

fromlen = sizeof(struct sockaddr_in);

while(1){
    n = recvfrom(sockfd, buf, 1024, 0, (struct sockaddr*)&from, &fromlen);
    if(n<0)
        error("Error on receiving from client");
    write(1, "Received a datagram: ", 21);
    write(1, buf, n);
    n = sendto(sockfd, "Message received\n", strlen("Message received\n"), 0, (struct sockaddr*)&from, fromlen);

    if(n<0)
        error("Error on sending to the client");
}

return 0;
}

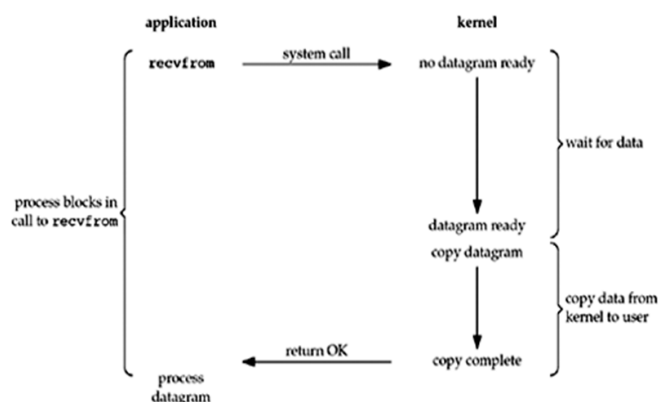
```

Advanced TCP/IP – I/O Modes

▼ Network operation steps - READ

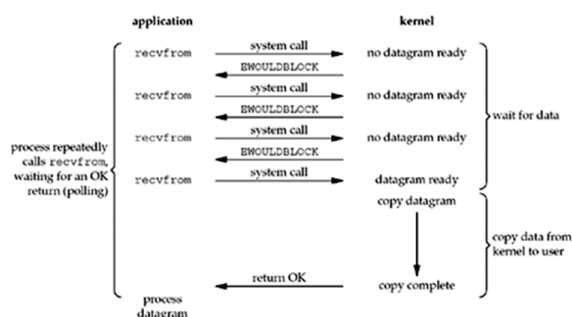
1. Wait for data to arrive from network.
2. Copying the data from the kernel to the process – i.e. copying the data from the kernel buffer into the application space.

1. Blocking I/O Model



- process blocks = process is sleeping
- process is sleeping while it waits for data
- when data arrives, the process wakes up and continues with its work

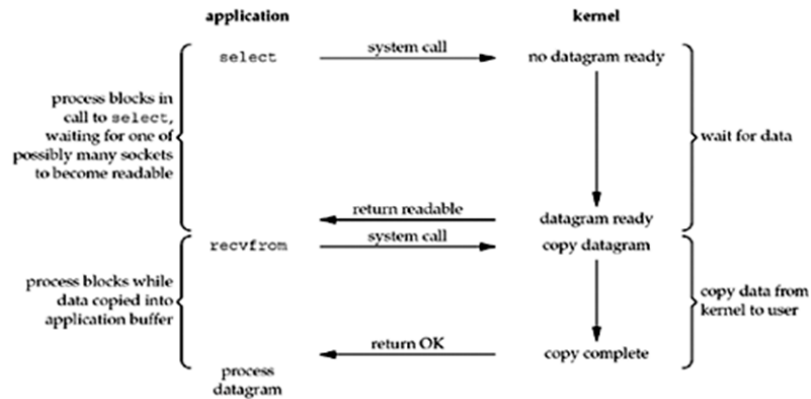
2. Non-Blocking I/O Model



- we can use `setsockopt` to set the socket to non blocking
- we are able to process other things between calls

3. Multiplexed I/O Model

- involves using select and poll



▼ select()

```
#include <sys/select.h>
#include <sys/time.h>

int select(int maxfd+1, fd_set* readset, fd_set* writeset, fd_set *exceptset,
const struct timeval *timeout)
//readset - watches sockets for read
//writeset - watches sockets for write
//timeout - how long?

void FD_ZERO(fd_set, *fdset); //clear all bits in fdset
void FD_SET(int fd, fd_set* fdset); //turn on the bit for fd in fdset
void FD_CLR(int fd, fd_set* fdset); //turn off the bit for fd in fdset
void FD_ISSET(int fd, fd_set* fdset); //is the fd ready?
```

- returns positive count of ready descriptors, 0 if timeout and -1 on error

Conditions for a socket to be ready

Condition	Readable	Writable	Exception
Data to read	Y		
Read half connection closed	Y		
New connection (for listen)	Y		
Space available for writing		Y	
Write half connection closed		Y	
Pending error	Y	Y	
TCP- out-of-bound DATA			Y